

# Coresets for Deep Learning

Hu Ding  
USTC

May 7, 2025

# Outline

- 1 Introduction
- 2 Coresets for Efficient Training of DL
- 3 Coresets for Improving Data Utilization of DL
- 4 Coresets for Large Language Models
- 5 Challenges

# Why Study Coreset for Deep Learning?

- **Data Explosion:** Modern deep learning models are trained on increasingly large datasets (e.g., ImageNet, OpenWebText), making training expensive and time-consuming.
- **Efficiency Bottleneck:** High computational and memory costs limit the scalability of deep models, especially on edge devices or under resource constraints.
- **Coreset Idea:** A *coreset* is a small, representative subset of the original dataset that preserves the model's training behavior.
- **Potential Benefits:**
  - ▶ Faster training and inference
  - ▶ Reduced storage and memory footprint
  - ▶ Improved generalization through better data selection
- **Challenge:** How to select informative and diverse samples that reflect both data distribution and learning dynamics?

# Coresets



Figure: The t-SNE visualization of coreset selection.

t-SNE<sup>1</sup>(t-Distributed Stochastic Neighbor Embedding) is a **nonlinear** dimensionality reduction algorithm commonly used to embed high-dimensional data into 2D or 3D space for visualization. Compared to **linear methods** like **PCA**, t-SNE is better at *capturing complex nonlinear structures* such as clusters.

<sup>1</sup>Van der Maaten L, Hinton G. Visualizing data using t-SNE. *Journal of machine learning research*, 2008, 9(11)

## 1 Introduction

## 2 Coresets for Efficient Training of DL

- Greedy Algorithms
- Example: CRAIG

## 3 Coresets for Improving Data Utilization of DL

- Active Learning
- Generative Models (GM)
- Continual Learning

## 4 Coresets for Large Language Models

- Coresets for Large Language Models

## 5 Challenges

- Challenges

# Greedy Algorithms

## Definition (Greedy Algorithm)

A **greedy algorithm** is any algorithm that follows the problem-solving **heuristic** of making the locally optimal choice at each stage<sup>a</sup>.

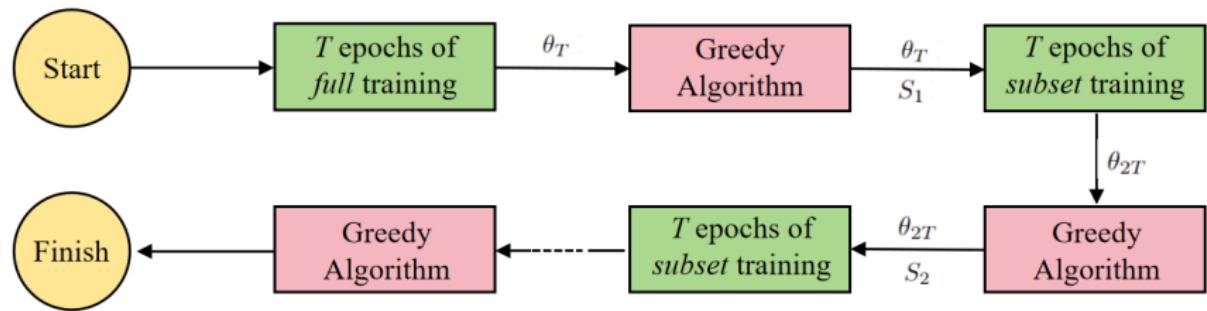
---

<sup>a</sup>Black, Paul E. (2 February 2005). “greedy algorithm”. *Dictionary of Algorithms and Data Structures*. U.S. National Institute of Standards and Technology (NIST). Retrieved 17 August 2012.

## Definition (Greedy Coreset Selection)

At each iteration (or epoch), the method appends the candidate samples that maximizes a **task-specific utility function**, thereby striving to preserve the essential statistical or geometric characteristics of the full dataset while substantially reducing its size.

# Greedy Coreset Selection in DL



**Figure:** Flowchart of greedy coresset selection in DL, where coresset selection is performed every  $T$  epochs and the model is trained on the selected coresset.

- ➊ Train on the entire dataset for  $T$  epochs (warm up).
- ➋ Based on the parameters obtained after the  $T$ -th epoch, perform a greedy search over the dataset to select a subset  $S$  of size  $k$  (*how to select?*).
- ➌ Train on the selected subset  $S$  for  $T$  epochs, then return to Step 2. Repeat until the algorithm converge.

# Foundations of Greedy Selection

- **Dataset.**

Let  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$  denote a dataset, where each  $(x_i, y_i)$  corresponds to an instance and its associated label or target.

- **Set Function.**

Consider a set function  $f : 2^{\mathcal{D}} \rightarrow \mathbb{R}$  that assigns a non-negative utility score  $f(\mathcal{S})$  to any subset  $\mathcal{S} \subseteq \mathcal{D}$ . The function  $f$  is assumed to reflect task-specific objectives.

- **Objective.**

The goal is to construct a representative subset  $\mathcal{S} \subseteq \mathcal{D}$  of fixed cardinality  $k$ , referred to as a *greedy coresset*, which approximately maximizes the utility function  $f(\mathcal{S})$ .

- **Marginal Gain.**

At each selection step, the algorithm identifies the element  $x \in \mathcal{D} \setminus \mathcal{S}$  that maximizes the marginal gain:

$$\Delta_f(x \mid \mathcal{S}) = f(\mathcal{S} \cup \{x\}) - f(\mathcal{S}),$$

where  $\Delta_f(x \mid \mathcal{S})$  quantifies the incremental benefit of adding  $x$  to the current subset  $\mathcal{S}$ .

# Algorithm Procedure

---

**Algorithm** Generic greedy coresset construction

---

**Input:** Dataset  $\mathcal{D}$ , utility function  $f$ , target size  $k$

- 1: Initialize  $S \leftarrow S_0$  /\* optional warm start \*/
  - 2: **while**  $|S| < k$  **do**
  - 3:    $x^* \leftarrow \arg \max_{x \in \mathcal{D} \setminus S} \Delta_f(x \mid S)$
  - 4:    $S \leftarrow S \cup \{x^*\}$
  - 5: **end while**
  - 6: **return**  $S$
-

# Cardinality Constrained Monotone Submodular Maximization (CCMSM)

## Definition (Submodular Function)

A set function  $f : 2^{\mathcal{D}} \rightarrow \mathbb{R}$ , defined over a finite ground set  $\mathcal{D}$ , is called submodular function if it satisfies

$$f(A \cup \{x\}) - f(A) \geq f(B \cup \{x\}) - f(B) \quad (1)$$

for all subsets  $A \subseteq B \subseteq \mathcal{D}$  and any element  $x \notin B$ . (*diminishing returns property*)

## Definition (CCMSM)

The cardinality-constrained monotone submodular maximization problem can be defined as follows:

$$\max_{S \subseteq \mathcal{D} : |S| \leq k} f(S). \quad (2)$$

## Theoretical guarantee

If  $f$  is *monotone submodular function* (i.e.,  $f(A) \leq f(B)$  whenever  $A \subseteq B$ ) and *normalized* (i.e.,  $f(\emptyset) = 0$ ), then the classical result by Nemhauser et al.<sup>2</sup> guarantees that Alg. 1 achieves the well-known  $(1 - 1/e)$  approximation ratio.

### Theorem (Nemhauser et al.)

Given a non-negative monotone submodular function  $f : 2^{\mathcal{D}} \rightarrow \mathbb{R}$  and the subsets  $\{S_i\}_{i \geq 0}$  extracted by Algorithm 1, we have, for all non-negative integers  $k, l$ :

$$f(S_l) \geq \left(1 - \exp\left(-\frac{l}{k}\right)\right) \max_{S: |S| \leq k} f(S), \quad (3)$$

and in particular, when  $l = k$ ,

$$f(S_k) \geq \left(1 - \frac{1}{e}\right) \max_{S: |S| \leq k} f(S). \quad (4)$$

---

<sup>2</sup>Nemhauser, George L., Laurence A. Wolsey, and Marshall L. Fisher. "An analysis of approximations for maximizing submodular set functions-I." *Mathematical programming* 14 (1978): 265-294.

## 1 Introduction

## 2 Coresets for Efficient Training of DL

- Greedy Algorithms
- Example: CRAIG

## 3 Coresets for Improving Data Utilization of DL

- Active Learning
- Generative Models (GM)
- Continual Learning

## 4 Coresets for Large Language Models

- Coresets for Large Language Models

## 5 Challenges

- Challenges

# CRAIG

---

## Coresets for Data-efficient Training of Machine Learning Models

---

Baharan Mirzasoleiman<sup>1</sup> Jeff Bilmes<sup>2</sup> Jure Leskovec<sup>3</sup>

**Figure:** <sup>1</sup>Department of Computer Science, University of California, Los Angeles, USA <sup>2</sup>Department of Electrical Engineering, University of Washington, Seattle, USA <sup>3</sup>Department of Computer Science, Stanford University, Stanford, USA. **Proceedings of the 37 th International Conference on Machine Learning, Vienna, Austria, PMLR 119, 2020.**

# Setting

## Training Machine Learning Models

Often reduces to minimizing a regularized empirical risk function

$$\begin{aligned}
 & \text{Feature Label} \\
 & \text{Training data volume: } \{(x_i, y_i), i \in V\} \\
 & w_* \in \arg \min_{w \in \mathcal{W}} f(w), \quad f(w) = \sum_{i \in V} f_i(w) + r(w), \quad f_i(w) = l(w, (x_i, y_i)) \\
 & \qquad \qquad \qquad \text{Regularizer} \\
 & \qquad \qquad \qquad \text{Loss function associated with} \\
 & \qquad \qquad \qquad \text{training example } i \in V
 \end{aligned}$$

Examples:

- Convex  $f(w)$ : Linear regression, logistic regression, ridge regression, regularized support vector machines (SVM)
- Non-convex  $f(w)$ : Neural networks

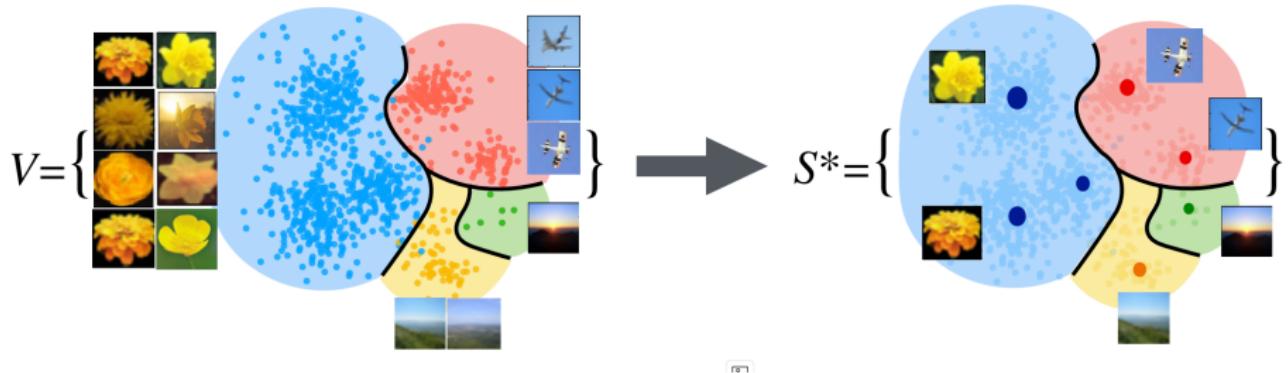
# Problem

## How to Find the “Right” Data for Machine Learning?

- The most informative subset

$$S^* = \arg \max_{S \subseteq V} F(S), \quad \text{s.t. } |S| \leq k$$

- What is a good choice for  $F(S)$ ?
- If we can find  $S^*$ , we get a approximate  $|V|/|S^*|$  speedup by only training on  $S^*$



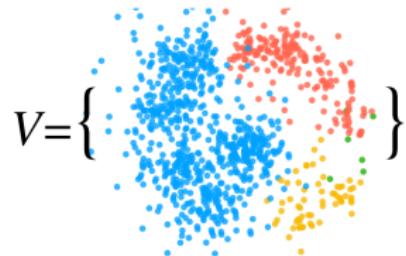
# Learning from Coresets

**Idea:** select the smallest subset  $S^*$  and weights  $\gamma$  that closely estimate the full gradient.

$$S^* = \arg \min_{S \subseteq V, \gamma_j \geq 0 \forall j} |S|, \quad \text{s.t.} \quad \max_{w \in \mathcal{W}} \left\| \sum_{i \in V} \nabla f_i(w) - \sum_{j \in S} \gamma_j \nabla f_j(w) \right\| \leq \epsilon$$

**Solution:** for every  $w \in \mathcal{W}$ ,  $S^*$  is the set of **exemplars** of all the data points in the **gradient space**.

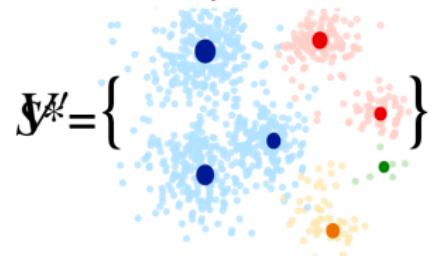
**Training Data:**  $\{(x_i, y_i), i \in V\}$



Gradients at  $w$

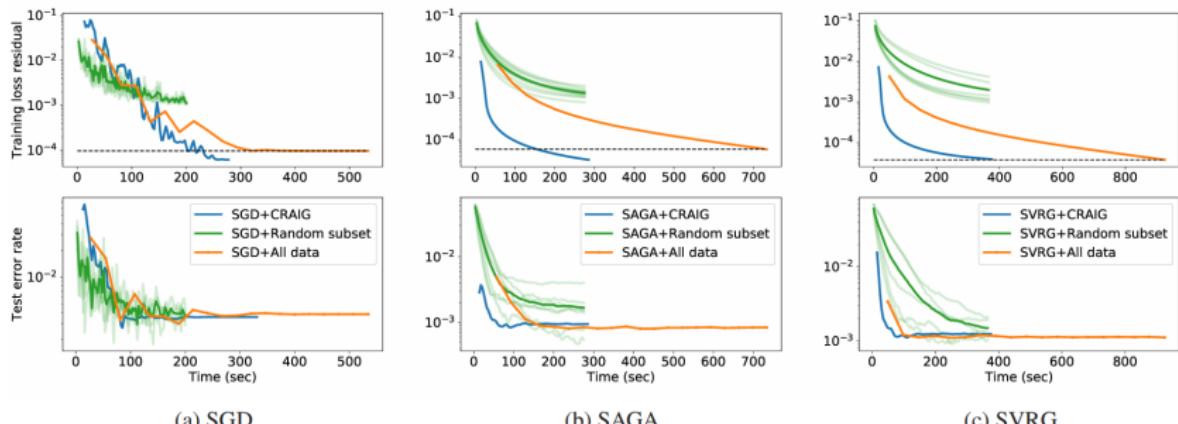


$V' = \{ \nabla f_i(w), i \in V \}$



# Application of CRAIG to Logistic Regression

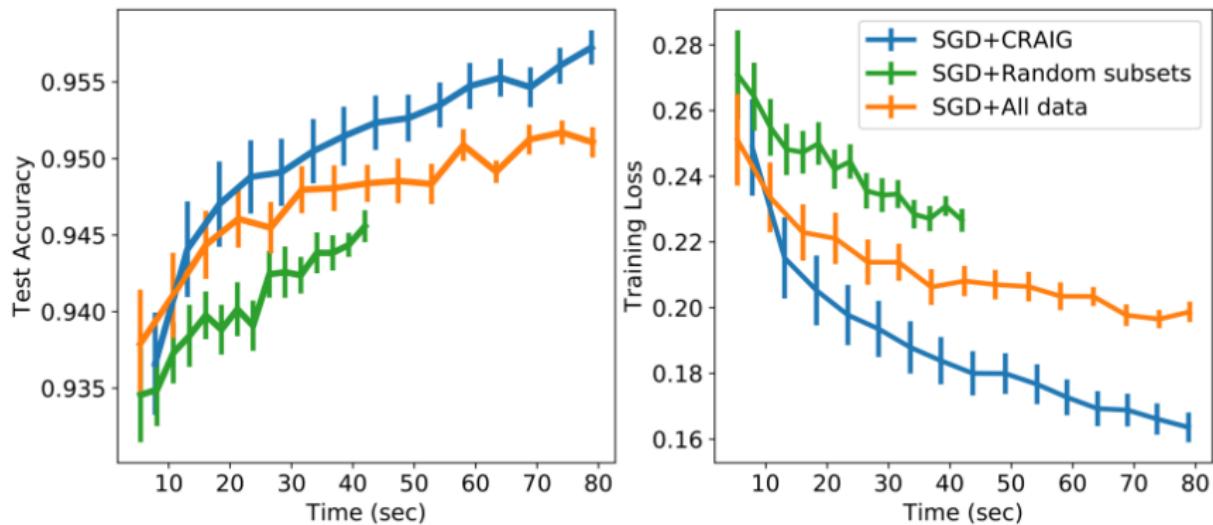
Training on subsets of size 10% of Covtype with 581K points



**Figure 1.** Loss residual and error rate of SGD, SVRG, SAGA for Logistic Regression on Covtype data set with 581,012 data points. We compare CRAIG (10% selected subset) (blue) vs. 10% random subset (green) vs. entire data set (orange). CRAIG gives the average speedup of 3x for achieving similar loss residual and error rate across the three optimization methods.

# Application of CRAIG to Neural Networks

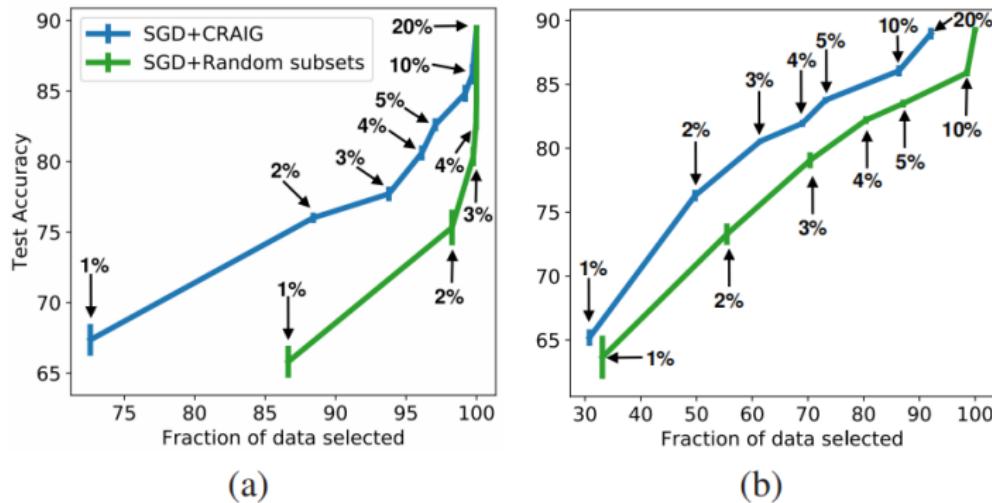
Training on MNIST with a 2-layer neural network with 50K points



**Figure:** Test accuracy and training loss of SGD applied to subsets found by CRAIG vs. random subsets on MNIST with a 2-layer neural network. CRAIG provides 2x to 3x speedup and a better generalization performance.

# Application of CRAIG to Deep Networks

Training ResNet20 on subsets of various size from CIFAR10 with 50K points



**Figure:** Test accuracy vs. fraction of data selected during training of ResNet20 on CIFAR10. (a) At the beginning of every epoch, a new subset of size 1%, 2%, 3%, 4%, 5%, 10%, or 20% is selected by CRAIG. (b) Every 5 epochs a new subset of similar size is selected by CRAIG. SGD is then applied to training on the selected subsets.

## 1 Introduction

## 2 Coresets for Efficient Training of DL

- Greedy Algorithms
- Example: CRAIG

## 3 Coresets for Improving Data Utilization of DL

- Active Learning
- Generative Models (GM)
- Continual Learning

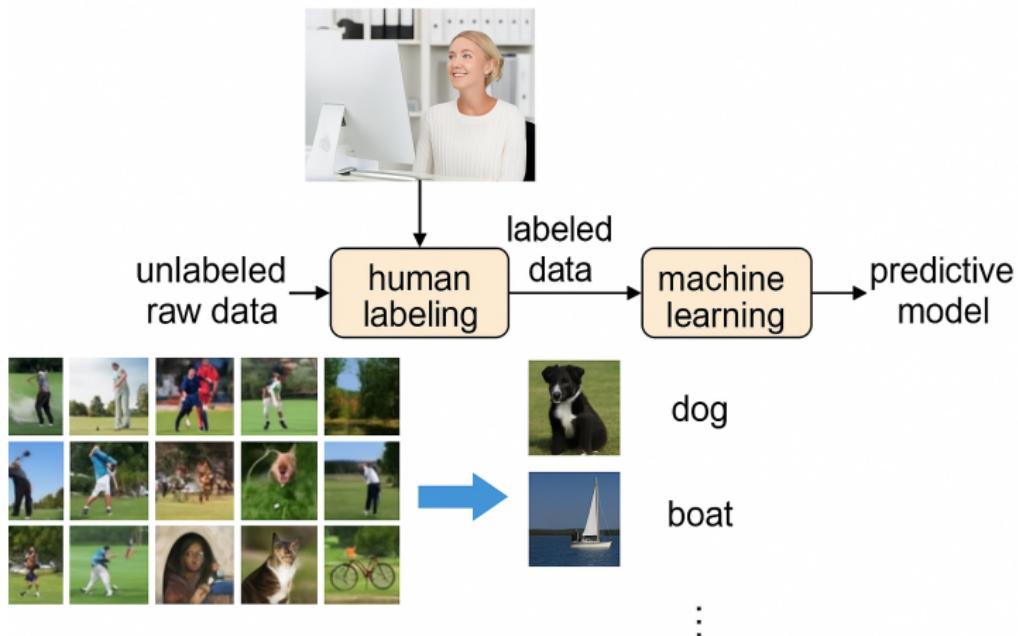
## 4 Coresets for Large Language Models

- Coresets for Large Language Models

## 5 Challenges

- Challenges

# Conventional (Passive) Machine Learning



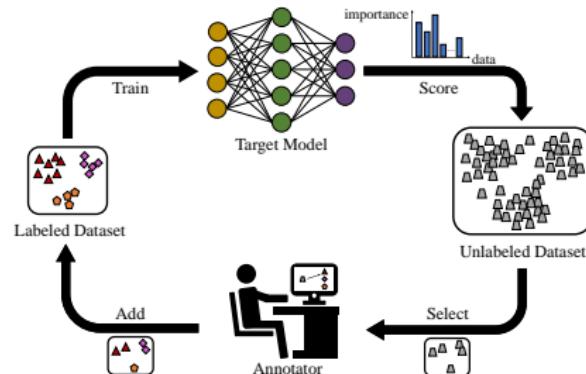
*Can we train machines with less labeled data and less human supervision?*

# Active Learning (Pool-based)

## Definition

**Active learning<sup>a</sup>** refers to the process of strategically selecting and labeling the most informative and representative samples from an unlabeled dataset to maximize learning efficiency and model accuracy.

<sup>a</sup>Settles, Burr. "Active learning literature survey." (2009).



**Figure:** Pool-based active learning repeatedly executes four key steps: (1) training the target model using the labeled data, (2) scoring the importance of unlabeled examples and selecting the most important ones, (3) obtaining annotations for the selected examples from annotators, and (4) incorporating the newly labeled data into the existing labeled dataset.

# Coreset Selection in AL

Based on the criteria for evaluating the importance of unlabeled examples, we categorize active learning approaches into three types: Loss-Based Methods (**Importance**); Coverage-Based Methods (**Diversity**); Hybrid Methods.

Type	Methods	Highlights
Loss-Based	CEAL [110]	Selecting examples using Least Confidence, Margin, or Entropy, querying true labels for some and assigning pseudo-labels to others.
	DFAL [162]	Selecting examples closest to decision boundaries.
	Yoo et al. [163]	Training a model to predict loss for unlabeled examples.
	Gao et al. [164]	Prioritizing examples with inconsistent predictions under data augmentations.
	SAAL [165]	Selecting examples exhibiting the highest training loss sharpness.
Coverage-Based	Sener et al. [114]	Formulating active learning as a $k$ -center problem and selecting examples farthest from the already selected ones.
	DAL [117]	Employing a binary classifier to discriminate between labeled and unlabeled examples.
	VAAL [118]	Adversarially training a discriminator to discriminate between labeled and unlabeled examples.
	ProbCover [138]	Formulating active learning as a Max Probability Cover problem.
	MaxHerding [139]	Selecting examples that maximizes their similarity to the entire unlabeled dataset.
Hybrid	BADGE [119]	Selecting examples by applying $k$ -means++ on their gradient embeddings.
	WAAL [127]	Formulating active learning as a distribution matching problem.
	NoiseStability [166]	Selecting examples based on the deviations in their feature embeddings under small perturbations.
	Cluster-Margin [120]	Performing hierarchical agglomerative clustering and selecting examples with low margin scores from the clusters.
	LDM-S [122]	Selecting high-diversity examples closest to decision boundaries.

# Coresets for AL

Published as a conference paper at ICLR 2018

---

## ACTIVE LEARNING FOR CONVOLUTIONAL NEURAL NETWORKS: A CORE-SET APPROACH

**Ozan Sener\***

Intel Labs

ozan.sener@intel.com

**Silvio Savarese**

Stanford University

ssilvio@stanford.edu

3

---

<sup>3</sup>Sener, Ozan, and Silvio Savarese. "Active Learning for Convolutional Neural Networks: A Core-Set Approach." International Conference on Learning Representations. 2018.

# Coresets for AL

## Coreset selection problem:

- Coreset selection problem aims to find a small subset given a large labeled dataset such that a model learned over the small subset is competitive over the whole dataset.

## Selects a batch of samples:

- choose “ $b$ ” center points such that the largest distance between a data point and its nearest center is minimized in feature space.

$$\min_{s: |s| \leq b} \max_i \min_{j \in s} \Delta(\mathbf{x}_i, \mathbf{x}_j)$$

---

### Algorithm 1 k-Center-Greedy

---

**Input:** data  $\mathbf{x}_i$ , existing pool  $s^0$  and a budget  $b$

Initialize  $s = s^0$

**repeat**

$u = \arg \max_{i \in [n] \setminus s} \min_{j \in s} \Delta(\mathbf{x}_i, \mathbf{x}_j)$

$s = s \cup \{u\}$

**until**  $|s| = b + |s^0|$

**return**  $s \setminus s^0$

---

# Experiments

Datasets: CIFAR-10, CIFAR-100, Caltech-256, SVHN

- ① The weakly-supervised model has access to labeled examples as well as unlabelled examples.
- ② The fully-supervised model only has access to the labeled data points.

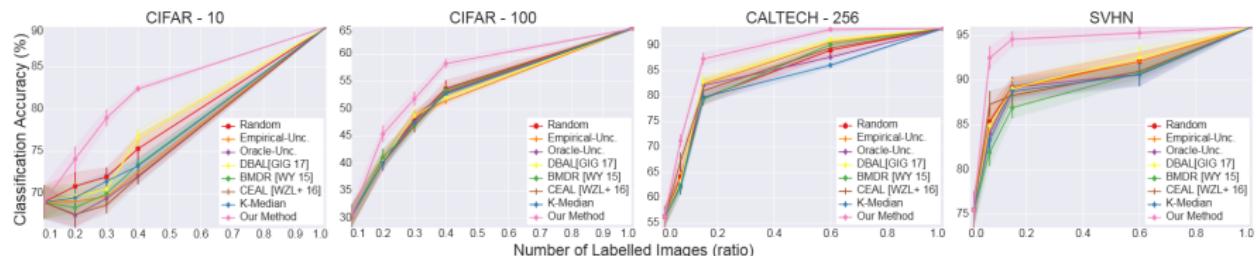


Figure 3: Results on Active Learning for Weakly-Supervised Model (error bars are std-dev)

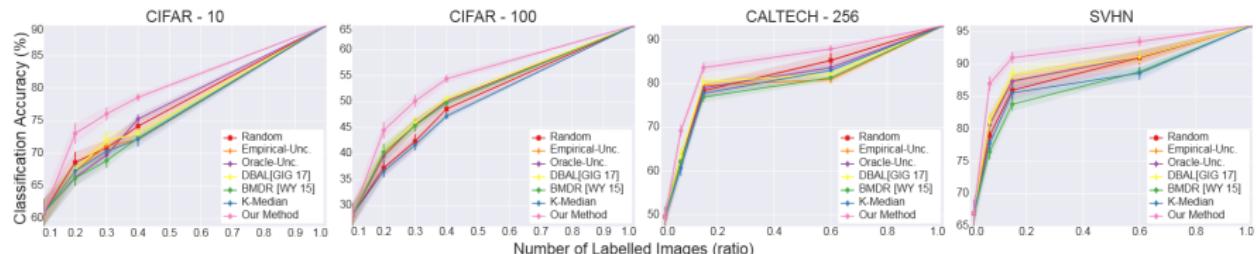


Figure 4: Results on Active Learning for Fully-Supervised Model (error bars are std-dev)

## 1 Introduction

## 2 Coresets for Efficient Training of DL

- Greedy Algorithms
- Example: CRAIG

## 3 Coresets for Improving Data Utilization of DL

- Active Learning
- Generative Models (GM)
- Continual Learning

## 4 Coresets for Large Language Models

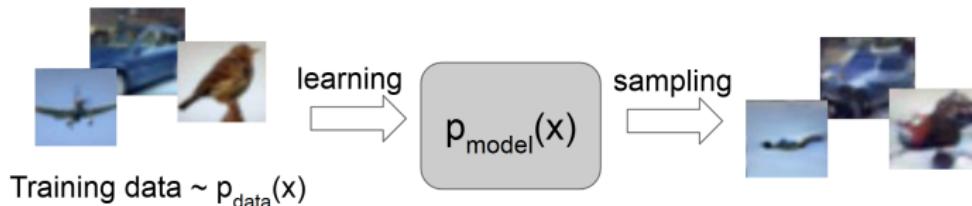
- Coresets for Large Language Models

## 5 Challenges

- Challenges

# Generative Models

Given training data, generate new samples from same distribution



## Objectives:

- ① Learn  $p_{\text{model}}(x)$  that approximates  $p_{\text{data}}(x)$
- ② Sampling new  $x$  from  $p_{\text{model}}(x)$

## Taxonomy of Generative Models

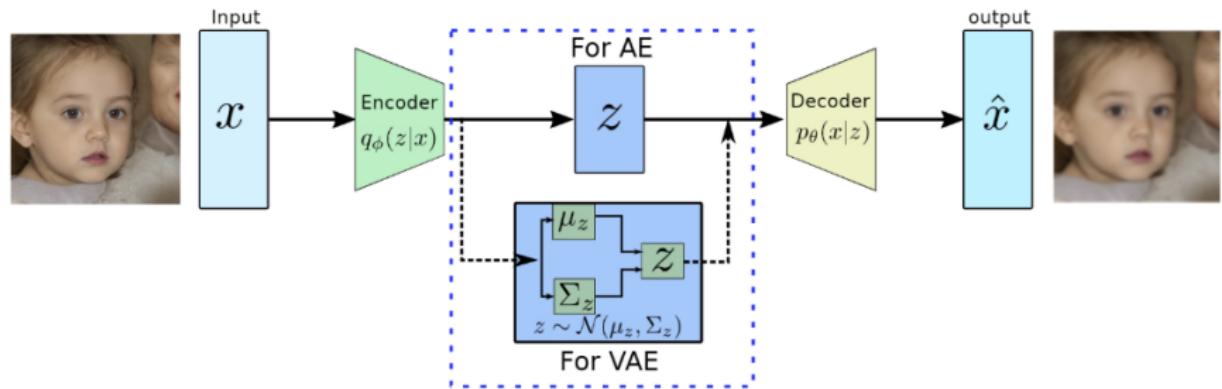
- ① Generative Adversarial Networks (GAN)
- ② Diffusion Models
- ③ Variational Autoencoders (VAE)

# Coreset selection in GM

Generative Model	Methods	Highlights
GANs	Small-GAN [136]	Formulating coreset selection as a $k$ -center problem with Euclidean distance as the distance metric.
	FDD-Coreset [137]	Formulating coreset selection as a $k$ -center problem with Fréchet distance as the distance metric.
	Top- $k$ GAN [108]	Discarding unrealistic synthetic images when training the generator.
	DeVries et al. [109]	Focusing on high-density regions from the data manifold.
	FIS-GAN [99]	Prioritizing challenging regions in the latent noise distribution.
VAEs	ByPE-VAE [126]	Employing Bayesian pseudocoresets to accelerate the training of Exemplar VAE.
Diffusion Models	Li et al. [142]	Coresets with learnable class-wise weights.
	Briq et al. [143]	Experimentally evaluating several coreset selection methods.

# Variational Autoencoder (VAE)

- VAE: a likelihood-based generative model



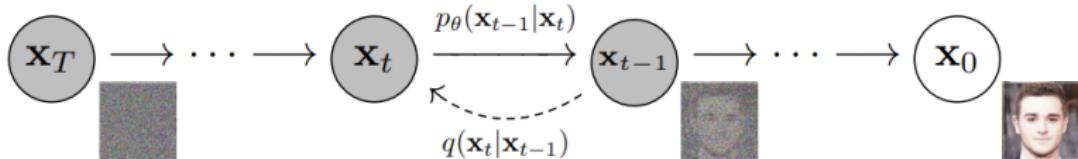
- Encoder: a model that approximates the posterior  $q(z|x)$
- Decoder: a model that transforms a Gaussian variable  $z$  to real data
- Training: maximize the ELBO

$$\mathcal{L}(\theta, \phi; x) = \mathbb{E}_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x) \parallel p_\theta(z)),$$

# Diffusion Model

## Denoising Diffusion Probabilistic Models (DDPM)<sup>4</sup>

- ➊ Diffusion models have two processes
- ➋ **Forward** diffusion process gradually adds noise to input
- ➌ **Reverse** denoising process learns to generate data by denoising



<sup>4</sup>Ho, Jonathan, Ajay Jain, and Pieter Abbeel. "Denoising diffusion probabilistic models." Advances in neural information processing systems 33 (2020): 6840-6851.

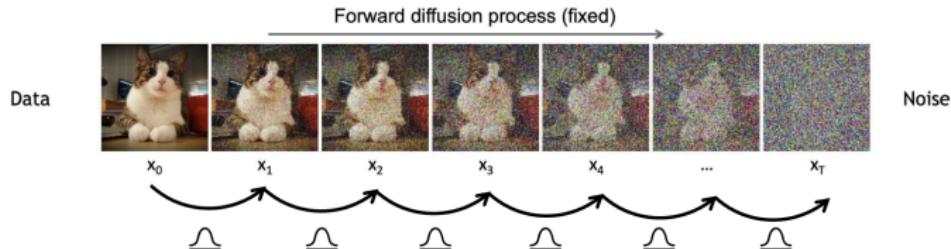
# Forward Diffusion Process

- ➊ Forward diffusion process is stacking **fixed VAE encoders**

- ▶ gradually adding **Gaussian noise** according to schedule  $\beta_t$

$$q(\mathbf{x}_t \mid \mathbf{x}_{t-1}) = \mathcal{N} \left( \mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I} \right)$$

$$q(\mathbf{x}_{1:T} \mid \mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t \mid \mathbf{x}_{t-1})$$



- ➋ The forward process allows sampling of  $\mathbf{x}_t$  at arbitrary timestep  $t$  in closed form:

$$q(\mathbf{x}_t \mid \mathbf{x}_0) = \mathcal{N} \left( \mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I} \right) \quad \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$$

- ➌ The noise schedule ( $\beta_t$  values) is designed such that  $q(\mathbf{x}_T \mid \mathbf{x}_0) \approx \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I})$

# Reverse Denoising Process

## ➊ Generation process

- ▶ Sample  $x_t \sim \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I})$
- ▶ Iteratively sample  $x_{t-1} \sim q(x_{t-1} | x_t)$

## ➋ $q(x_{t-1} | x_t)$ not directly tractable

➌ But can be estimated with a Gaussian distribution if  $\beta_t$  is small at each step

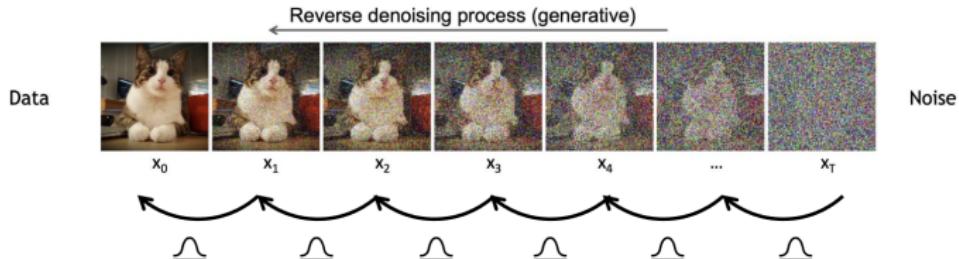
## ➍ Reverse diffusion process is stacking learnable VAE decoders

- ▶ Predicting the mean and std of added Gaussian Noise

$$p(\mathbf{x}_T) = \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I}) \quad p_\theta(\mathbf{x}_{0:T}) = p(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$$

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \underbrace{\mu_\theta(\mathbf{x}_t, t), \sigma_t^2 \mathbf{I}}_{\text{Trainable Network, Shared Across All Timesteps}})$$

Trainable Network, Shared Across All Timesteps



# Learning the Denoising Model

- ① Denoising models are trained with variational upper bound (negative ELBO), as VAEs

$$\mathbb{E}_{q(\mathbf{x}_0)} [-\log p_\theta(\mathbf{x}_0)] \leq \mathbb{E}_{q(\mathbf{x}_0)q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ -\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] =: \mathcal{L}$$

- ② In DDPM, this is further simplified to:

$$L = \mathbb{E}_{\mathbf{x}_0 \sim q(\mathbf{x}_0), \epsilon \sim \mathcal{N}(0, \mathbf{I}), t \sim \mathcal{U}(1, T)} \left[ \left\| \epsilon - \epsilon_\theta \left( \underbrace{\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon}_{\mathbf{x}_t}, t \right) \right\|^2 \right]$$

- ③ Summary: Training and Sampling

---

### Algorithm 1 Training

---

```

1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
      
$$\nabla_\theta \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2$$

6: until converged

```

---

### Algorithm 2 Sampling

---

```

1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:   
$$\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$$

5: end for
6: return  $\mathbf{x}_0$ 

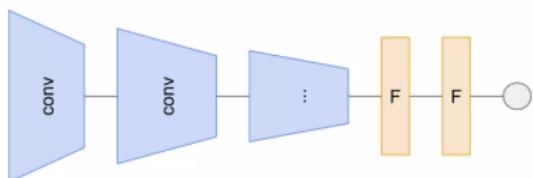
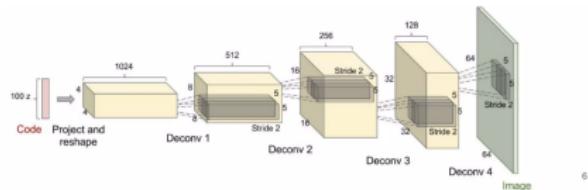
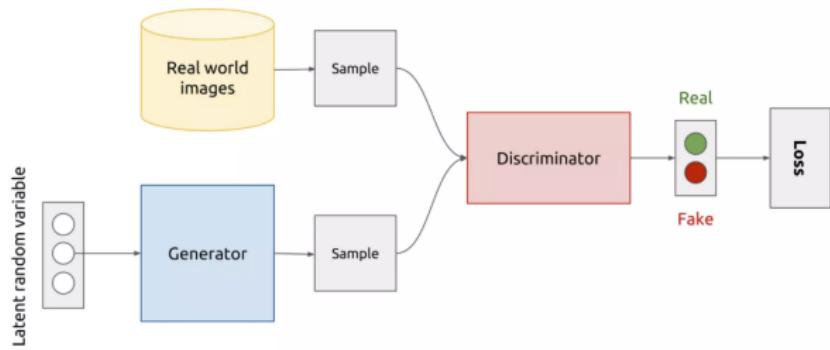
```

# Generative Adversarial Networks (GANs)

- A method of training deep generative models
- **Idea:** A **generator** and a **discriminator** against each other
- Generator tries to draw samples from  $p_{\text{model}}(x)$
- Discriminator tries to tell if sample came from the generator or the real world
- Both discriminator and generator are deep networks (differentiable functions)
- Can train with backprop: train discriminator for a while, then train generator, then discriminator, . . .

# Generative Adversarial Networks (GAN)

## GAN's Architecture<sup>5</sup>



<sup>5</sup> [slideshare.net/slideshow/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016/64667744](https://slideshare.net/slideshow/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016/64667744)

# Small-GAN: Speeding Up GAN Training Using Core-sets<sup>6</sup>

- ① Training with very large mini-batches has been shown to significantly improve performance [Brock et al. (2018)].
- ② But using large batches is slow and expensive.

Batch-Sizes	256	512	1024	2048
FID	18.65	15.30	14.88	12.39

How can we use large batch-size without overhead cost?

If we could generate batches that were effectively large though actually small.

---

<sup>6</sup>Sinha, Samarth, et al. "Small-gan: Speeding up gan training using core-sets." International Conference on Machine Learning. PMLR, 2020.

## GreedyCoreset

**Setting:** In computational geometry, a Coreset  $Q$  of a set  $P$  is a subset  $Q \subset P$  that approximates the ‘shape’ of  $P$  [Agarwal et al., 2005].

The Coreset selection problem can be formulated as the *minimax facility location* formulation:

$$\min_{Q: |Q|=k} \max_{x_i \in P} \min_{x_j \in Q} d(x_i, x_j)$$

- Exact solution to set-covers is **NP-Hard**
- We can use Greedy Approximations

### Algorithm 1 GreedyCoreset

**Input:** batch size ( $k$ ), data points ( $x$  where  $|x| > k$ )

**Output:** subset of  $x$  of size  $k$

```

 $s \leftarrow \{\}$                                  $\triangleright$  Initialize the sampled set
 $\text{while } |s| < k \text{ do}$                    $\triangleright$  Iteratively add points to sampled set
     $p \leftarrow \arg \max_{x_i \notin s} \min_{x_j \in s} d(x_i, x_j)$ 
     $s \leftarrow s \cup \{p\}$ 
 $\text{end while}$   $\text{return } s$ 

```

# Small-GAN

---

**Algorithm 2** Small-GAN
 

---

**Input:** target batch size ( $k$ ), starting batch size ( $n > k$ ), Inception embeddings ( $\phi_I$ )

**Output:** a trained GAN

Initialize networks  $G$  and  $D$

**for**  $step = 1$  to ... **do**

$z \sim p(z)$

$x \sim p(x)$

$\phi(x) \leftarrow \phi_I(x)$

$\hat{z} \leftarrow \text{GreedyCoreset}(z)$

$\widehat{\phi(x)} \leftarrow \text{GreedyCoreset}(\phi(x))$

$\hat{x} \leftarrow \phi_I^{-1}(\widehat{\phi(x)})$

    Update GAN parameters as usual

**end for**

---

▷ Sample  $n$  points from the prior

▷ Sample  $n$  points from the data distribution

▷ Get cached embeddings for  $x$

▷ Get Core-set of  $z$

▷ Get Core-set of embeddings

▷ Get  $x$  corresponding to sampled embeddings

**Q:** Why use Inception embeddings  $\phi_I(\cdot)$ ?

# Small-GAN

---

**Algorithm 2** Small-GAN
 

---

**Input:** target batch size ( $k$ ), starting batch size ( $n > k$ ), Inception embeddings ( $\phi_I$ )

**Output:** a trained GAN

Initialize networks  $G$  and  $D$

**for**  $step = 1$  to ... **do**

$z \sim p(z)$

$x \sim p(x)$

$\phi(x) \leftarrow \phi_I(x)$

$\hat{z} \leftarrow \text{GreedyCoreset}(z)$

$\widehat{\phi(x)} \leftarrow \text{GreedyCoreset}(\phi(x))$

$\hat{x} \leftarrow \phi_I^{-1}(\widehat{\phi(x)})$

    Update GAN parameters as usual

**end for**

---

▷ Sample  $n$  points from the prior

▷ Sample  $n$  points from the data distribution

▷ Get cached embeddings for  $x$

▷ Get Core-set of  $z$

▷ Get Core-set of embeddings

▷ Get  $x$  corresponding to sampled embeddings

**Q:** Why use Inception embeddings  $\phi_I(\cdot)$ ?

**A:** Running GreedyCoreset on images is ineffective because of noise and high dimensionality [Donoho et al., 2000]. Inception embeddings of images are cheap and effective

## Experiments:CIFAR10

**Small-GAN significantly outperforms regular GAN training at all batch-sizes**

GAN (batch-size = 128)	Small-GAN (batch-size = 128)	GAN (batch-size = 256)	Small-GAN (batch-size = 256)	GAN (batch-size = 512)	Small-GAN (batch-size = 512)
$18.75 \pm 0.2$	<b><math>16.73 \pm 0.1</math></b>	$17.9 \pm 0.1$	<b><math>16.22 \pm 0.3</math></b>	$15.68 \pm 0.2$	<b><math>15.08 \pm 0.1</math></b>

**Figure:** FID scores for CIFAR using SN-GAN

## Experiments:LSUN+ImageNet

**Small-GAN able to scale well to large scale settings**

<b>Small-GAN (64)</b>	<b>GAN (64)</b>	<b>GAN (128)</b>	<b>GAN (256)</b>
13.08	14.82	13.02	12.63

**Table:** FID scores using SAGAN for LSUN dataset

<b>GAN</b>	<b>Small-GAN</b>
19.40	<b>17.33</b>

**Table:** FID scores using SAGAN for ImageNet dataset

## 1 Introduction

## 2 Coresets for Efficient Training of DL

- Greedy Algorithms
- Example: CRAIG

## 3 Coresets for Improving Data Utilization of DL

- Active Learning
- Generative Models (GM)
- Continual Learning

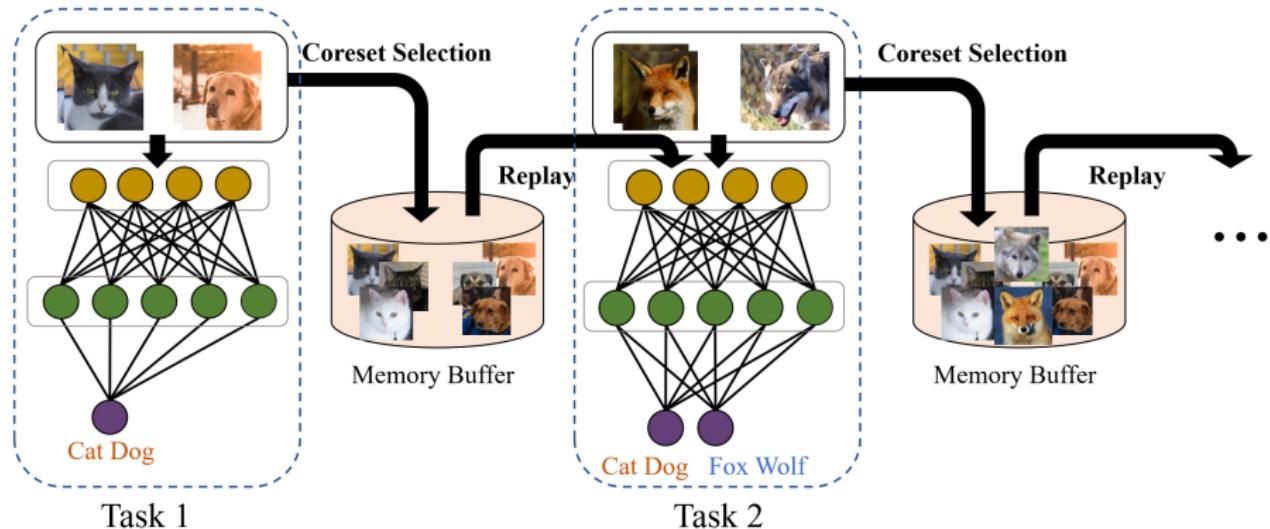
## 4 Coresets for Large Language Models

- Coresets for Large Language Models

## 5 Challenges

- Challenges

# Continual Learning (Experience rehearsal based or Coreset based)



**Figure:** Overview of the coreset-based method for continual learning. This approach aims to approximate and recover previous data distributions by storing a small subset of past training samples in a limited *memory buffer*. By leveraging this stored subset, the method helps mitigate *catastrophic forgetting*, ensuring that the model retains knowledge from earlier learning phases while adapting to new data.

# Coresets in CL

**Table 3** Coreset selection methods in continual learning.

Type	Methods	Highlights
Gradient-Free Strategies	Reservoir Sampling [193, 194, 195, 196]	A fundamental streaming algorithm for maintaining fixed-size uniform samples from unbounded data streams of unknown cardinality.
	Ring Buffer [197, 198, 199]	A class-partitioned memory architecture designed to ensure balanced exemplar retention.
	Herding [200, 201]	A method that dynamically adjusts slice dimensions by allocating $ \mathcal{M} /C_{seen}$ slots per class based on the number of observed classes $C_{seen}$ .
	$k$ -Means [134]	Use online $k$ -Means to estimate the $k$ class centroids in the penultimate feature space and store the input samples closest to each centroid in memory.
Gradient-Driven Strategies	GSS [202]	Formulate coreset selection as a constraint optimization problem of continual learning.
	OCS [203]	A coreset selection method that selects the most representative and informative cores at each iteration and trains them in an online manner.
	Greedy Coreset [63]	Propose a coreset construction via cardinality-constrained bilevel optimization.
	PBCS [82]	A continuous probabilistic bilevel formulation of coreset selection by learning a probabilistic weight for each training sample.
	BCSR [81]	A new bilevel formulation where the inner problem minimizes expected training error from a sampled distribution, and the outer problem learns a sparse distribution with $K$ nonzero entries to minimize overall training error.
	GCR [192]	A method for selecting a coreset that approximates the model parameter gradients over all previously seen data.
	RM [204]	A method that emphasize the importance of maintaining sample diversity in coreset.
	ASER [205]	Propose a novel Adversarial Shapley value scoring method for coreset selection in continual learning.

# Gradient-free strategies

## ① Reservoir Sampling

---

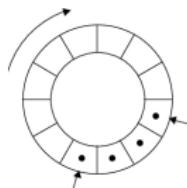
### Algorithm 3 - Reservoir Sampling

---

```
Input: memory buffer  $\mathcal{M}$ , number of seen examples  $N$ , example  $x$ , label  $y$ .  
if  $\mathcal{M} > N$  then  
     $\mathcal{M}[N] \leftarrow (x, y)$   
else  
     $j = \text{randomInteger}(\min = 0, \max = N)$   
    if  $j < |\mathcal{M}|$  then  
         $\mathcal{M}[j] \leftarrow (x, y)$   
    end if  
end if  
return  $M$ 
```

---

## ② Ring Buffer



## ③ Herding

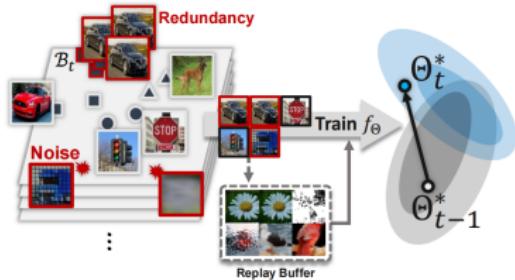
## ④ k-Means

# Gradient-driven strategies

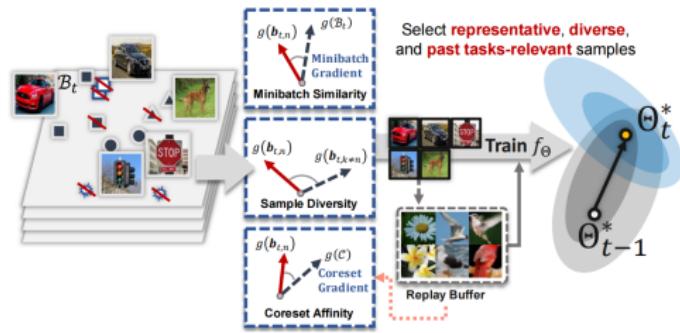
## ONLINE CORESET SELECTION FOR REHEARSAL-BASED CONTINUAL LEARNING

**Jaehong Yoon<sup>1</sup> Divyam Madaan<sup>2\*</sup> Eunho Yang<sup>1,3</sup> Sung Ju Hwang<sup>1,3</sup>**  
KAIST<sup>1</sup> New York University<sup>2</sup> AITRICS<sup>3</sup>  
jaehong.yoon@kaist.ac.kr, divyam.madaan@nyu.edu,  
eunhoy@kaist.ac.kr, sjhwang82@kaist.ac.kr

# OCS



(a) Existing rehearsal-based CL



(b) Online Coreset Selection (OCS)

**Figure:** Illustration of existing rehearsal-based CL and Online Coreset Selection (OCS).

- Existing rehearsal-based methods train on all the arrived instances and memorize a fraction of them in the replay buffer, which results in a suboptimal performance due to the **outliers (noisy or biased instances)**.
- OCS obtains the coresets by leveraging our **three selection strategies**, which discard the outliers at each iteration.

# Three Different Selection Strategies

Let  $b_{t,n} = \{x_{t,n}, y_{t,n}\} \in \mathcal{B}_t$  denote the  $n$ -th pair of data point with gradient  $\nabla f_\Theta(\mathbf{b}_{t,n})$  and its corresponding label at task  $\mathcal{T}_t$ .

## ① Minibatch similarity

$$S(b_{t,n} | \mathcal{B}_t) = \frac{\nabla f_\Theta(b_{t,n}) \bar{\nabla} f_\Theta(\mathcal{B}_t)^\top}{\|\nabla f_\Theta(b_{t,n})\| \cdot \|\bar{\nabla} f_\Theta(\mathcal{B}_t)\|}.$$

## ② Sample diversity

$$\mathcal{V}(b_{t,n} | \mathcal{B}_t \setminus b_{t,n}) = \frac{-1}{N_t - 1} \sum_{p \neq n}^{N_t-1} \frac{\nabla f_\Theta(b_{t,n}) \nabla f_\Theta(b_{t,p})^\top}{\|\nabla f_\Theta(b_{t,n})\| \cdot \|\nabla f_\Theta(b_{t,p})\|}.$$

## ③ Coreset affinity

$$\mathcal{A}(b_{t,n} | \mathcal{B}_C \sim C) = \frac{\nabla f_\Theta(b_{t,n}) \bar{\nabla} f_\Theta(\mathcal{B}_C)^\top}{\|\nabla f_\Theta(b_{t,n})\| \cdot \|\bar{\nabla} f_\Theta(\mathcal{B}_C)\|}.$$

# Algorithm

---

**Algorithm 1** Online Coreset Selection (OCS)
 

---

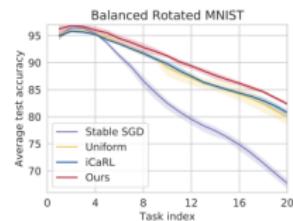
**input** Dataset  $\{\mathcal{D}_t\}_{t=1}^T$ , neural network  $f_\Theta$ , hyperparameters  $\lambda, \tau$ , replay buffer  $\mathcal{C} \leftarrow \{\}$ , buffer size bound  $J$ .

- 1: **for** task  $\mathcal{T}_t = \mathcal{T}_1, \dots, \mathcal{T}_T$  **do**
- 2:    $\mathcal{C}_t \leftarrow \{\}$  ▷ Initialize coreset for current task
- 3:   **for** batch  $\mathcal{B}_t \sim \mathcal{D}_t$  **do**
- 4:      $\mathcal{B}_C \leftarrow \text{SAMPLE}(\mathcal{C})$  ▷ Randomly sample a batch from the replay buffer
- 5:      $\mathbf{u}^* = \underset{n \in \{0, \dots, |\mathcal{B}_t| - 1\}}{\operatorname{argmax}} (\kappa) \mathcal{S}(\mathbf{b}_{t,n} \mid \mathcal{B}_t) + \mathcal{V}(\mathbf{b}_{t,n} \mid \mathcal{B}_{t \setminus b_{t,n}}) + \tau \mathcal{A}(\mathbf{b}_{t,n} \mid \mathcal{B}_C)$  ▷ Coreset selection
- 6:      $\Theta \leftarrow \Theta - \eta \nabla f_\Theta(\mathcal{B}_t[\mathbf{u}^*] \cup \mathcal{B}_C)$  with Equation (9) ▷ Model update with selected instances
- 7:      $\mathcal{C}_t \leftarrow \mathcal{C}_t \cup \widehat{\mathcal{B}}_t$
- 8:   **end for**
- 9:    $\mathcal{C} \leftarrow \mathcal{C} \cup \text{SELECT}(\mathcal{C}_t, \text{size} = J/T)$  with Equation (8) ▷ Memorize coreset in the replay buffer
- 10: **end for**

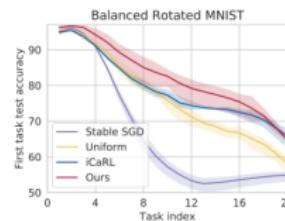
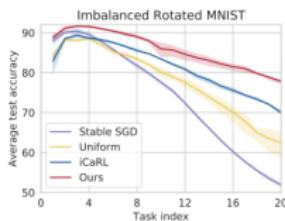
---

# Experiments

The imbalanced setting contains a different number of training examples for each class in a task.



(a) Average test accuracy



(b) First task test accuracy

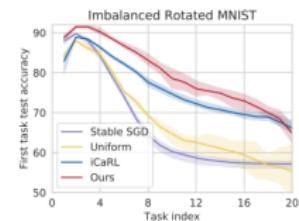
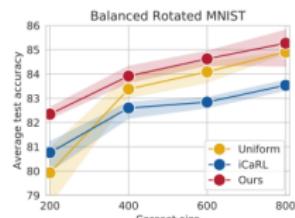
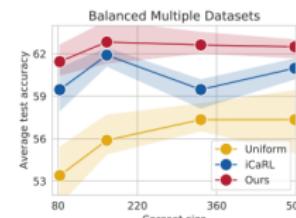
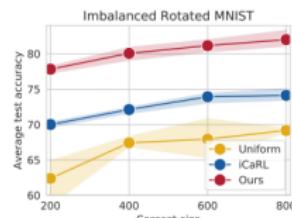


Figure: (a) Average accuracy (b) First task accuracy for balanced/imbalanced Rotated MNIST during CL.



(a) Rotated MNIST



(b) Multiple Datasets

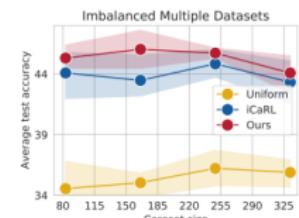


Figure: Performance comparison on various coreset sizes for balanced/imbalanced continual learning

## 1 Introduction

## 2 Coresets for Efficient Training of DL

- Greedy Algorithms
- Example: CRAIG

## 3 Coresets for Improving Data Utilization of DL

- Active Learning
- Generative Models (GM)
- Continual Learning

## 4 Coresets for Large Language Models

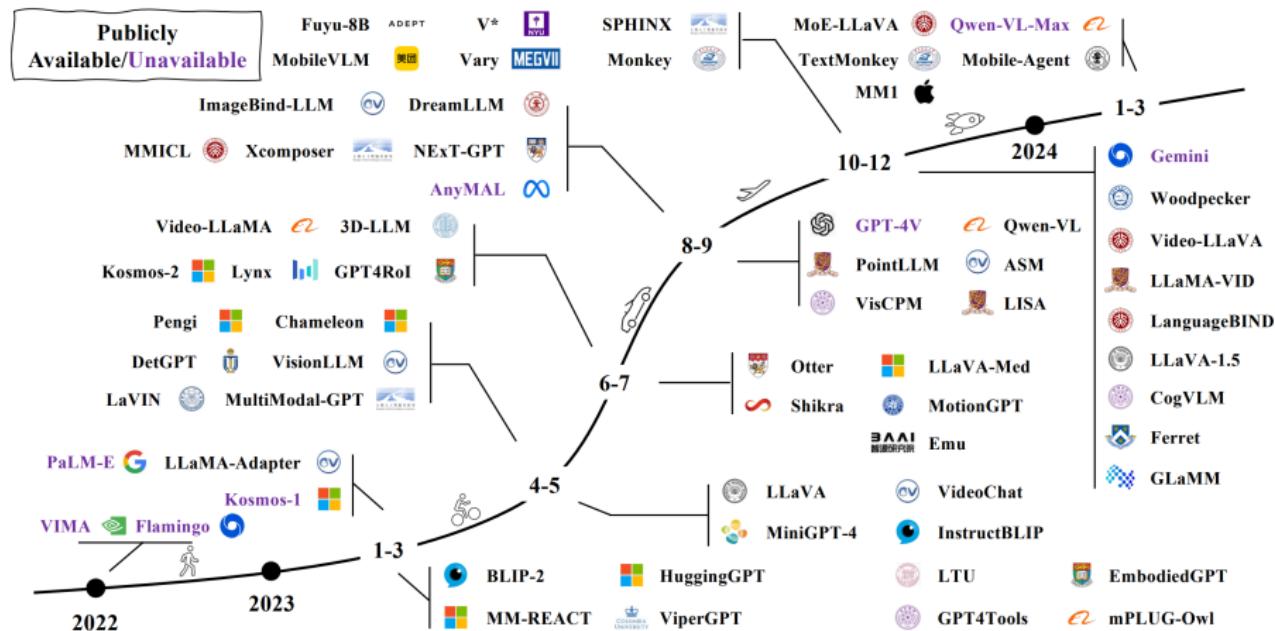
- Coresets for Large Language Models

## 5 Challenges

- Challenges

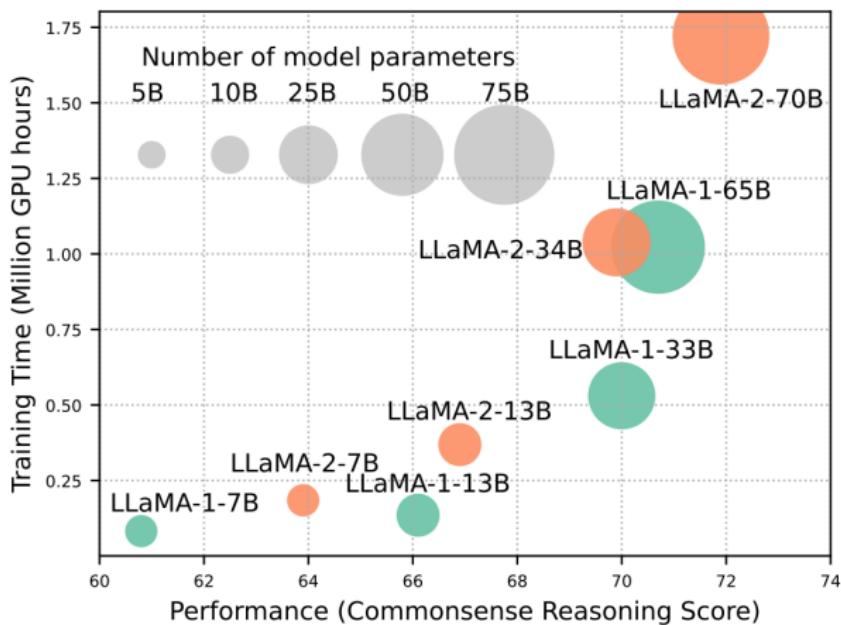
## Introduction

**Large Language Models (LLMs) are a type of advanced AI models designed to understand and generate human languages.**



## Efficient LLMs

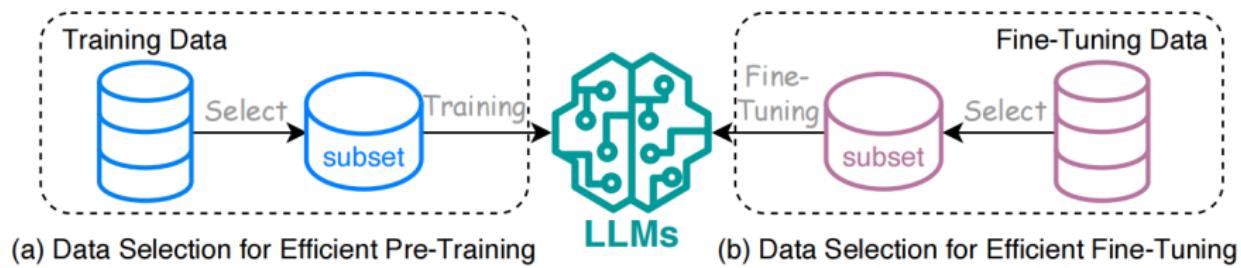
Although LLMs are leading the next wave of AI revolution, their remarkable capabilities come at substantial resource demands



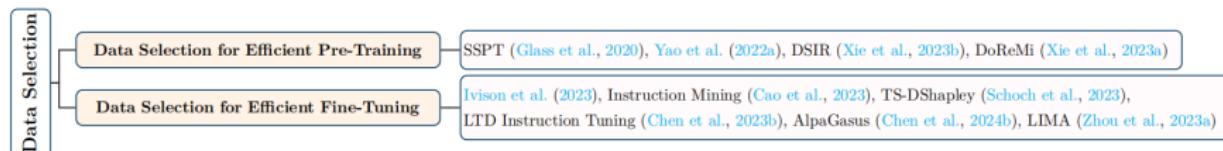
**Figure:** Illustration of model performance and model training time in GPU hours of LLaMA models at different scales.

# Coresets in LLMs

Coreset selection is a fundamental technique for enhancing efficiency. As summarized in Figure below, in the context of LLMs, coreset selection techniques have been primarily used for enhancing the efficiency of pretraining and fine-tuning.



**Figure:** Illustrations of coreset selection techniques for LLMs.



**Figure:** Summary of coreset selection techniques for LLMs.

Published as a conference paper at ICLR 2025

---

# MINI-BATCH CORESETS FOR MEMORY-EFFICIENT LANGUAGE MODEL TRAINING ON DATA MIXTURES

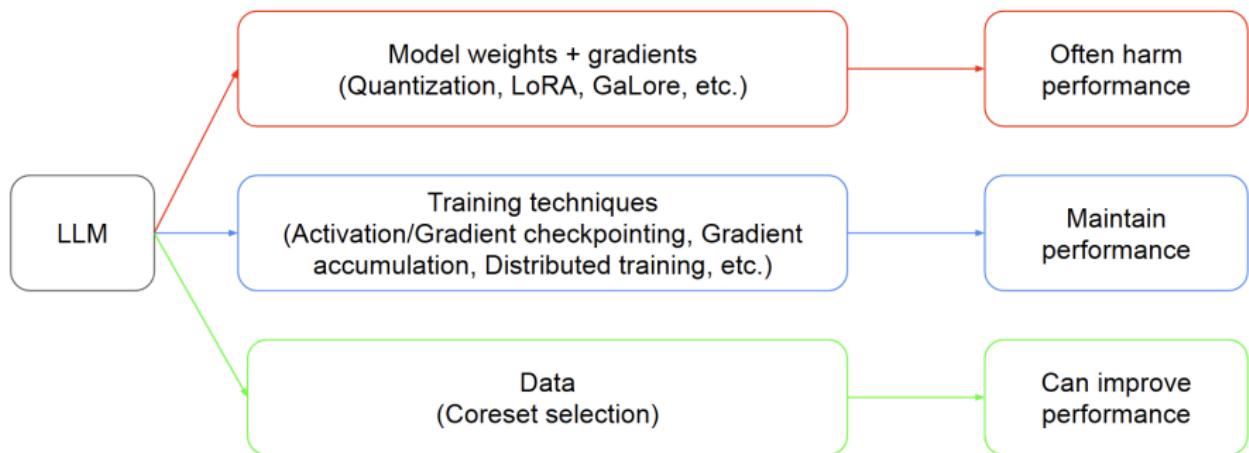
**Dang Nguyen    Wenhan Yang    Rathul Anand    Yu Yang    Baharan Mirzasoleiman**

{dangnth, hangeryang18, rathul, yuyang, baharan}@cs.ucla.edu

Computer Science Department, UCLA

# CoLM

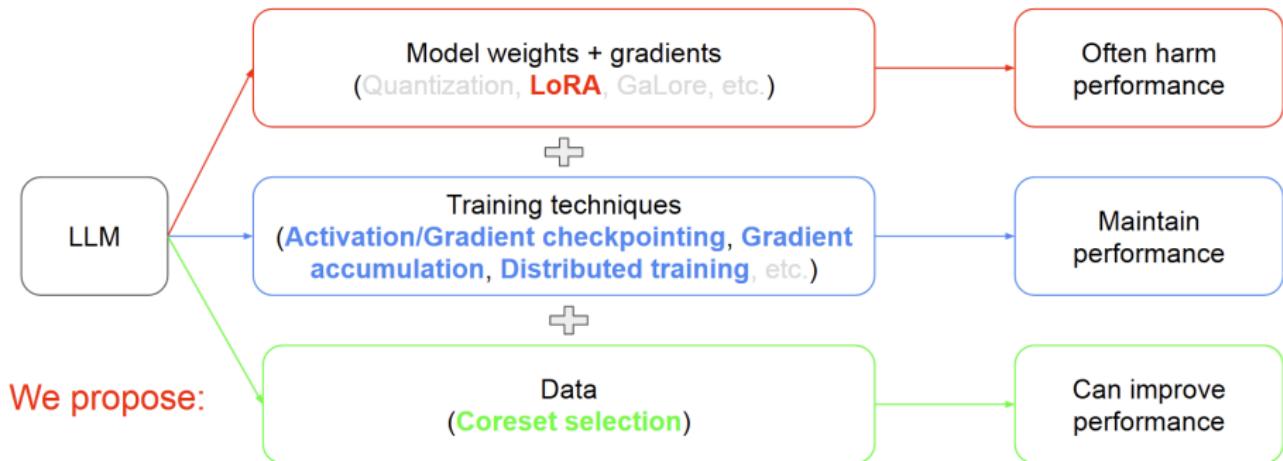
## How to reduce GPU memory requirement in LLM?



# CoLM

## How to reduce GPU memory requirement in LLM?

It's possible to stack three different approaches together.



# CoLM

## (Mini-batch) coresnet selection

- Training with larger mini-batch has a small variance, thus converges faster
- A mini-batch coresnet is a subset that has the similar gradient to the original mini-batch.  
Thus, it can be found by solving the gradient matching problem

$$S_t^* \in \arg \max_{S \subset \mathcal{M}_t^L, |S| \leq b} \sum_{i \in \mathcal{M}_t^L} \max_{s \in S} [C - \|g_{i,t} - g_{s,t}\|],$$

Coreset size                      Big constant

Minibatch coresnet      Original minibatch

Gradient of original minibatch      Gradient of minibatch coresnet

# CoLM

## Challenges of coreset selection for training LLMs

- ① Highly Imbalanced Language Data
- ② Adam optimizer
- ③ Very Large Gradient Dimensionality

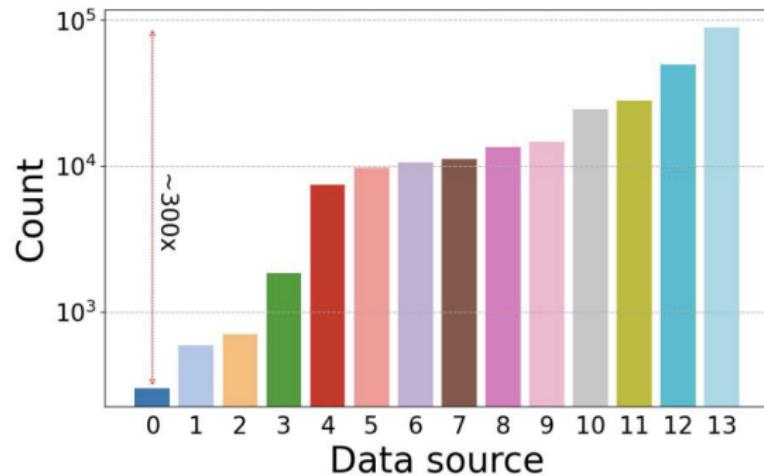


Figure: The number of samples from different data sources in MathInstruct.

# CoLM

## Coresets for Training LLMs (CoLM)

- ➊ A balanced sampling strategy
- ➋ Adam-like gradient normalization
- ➌ Sparsified zeroth-order gradient estimation

# Experimental results

CoLM yields the best of both worlds, increasing accuracy while reducing time & memory consumption.

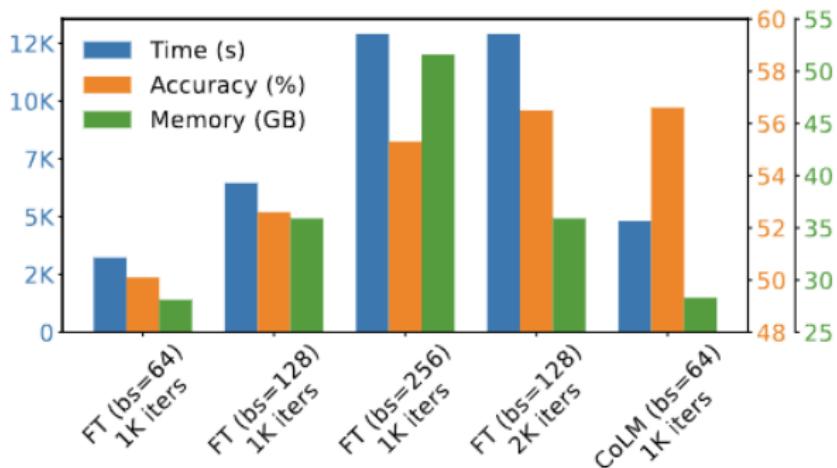
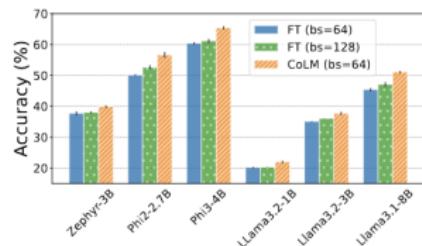


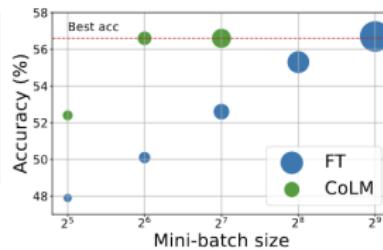
Figure: Time vs Accuracy vs Memory of fine-tuning Phi-2 with LoRA on MathInstruct.

# Experimental results

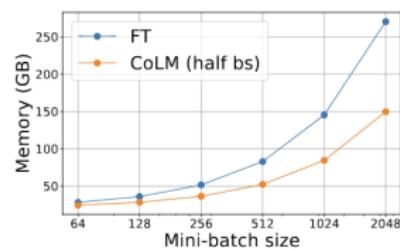
**CoLM improves the performance of different models and batch sizes.**



(a) Different LLMs



(b) Changing mini-batch size



(c) Memory usage vs batch size

**Figure:** a) CoLM with  $bs = 64$  (from 128) outperforms fine-tuning different models with  $bs = 64$  and  $bs = 128$  by a large margin; (b) CoLM improves the performance of training with different batch sizes. The size of each circle is proportional to the training time of the corresponding method. (c) CoLM reduces memory consumption, with reduction increasing as the batch size grows.

## 1 Introduction

## 2 Coresets for Efficient Training of DL

- Greedy Algorithms
- Example: CRAIG

## 3 Coresets for Improving Data Utilization of DL

- Active Learning
- Generative Models (GM)
- Continual Learning

## 4 Coresets for Large Language Models

- Coresets for Large Language Models

## 5 Challenges

- Challenges

# Challenges

- ① Dynamic and Adaptive Coreset Selection
- ② Automated and Multi-objective Coreset Optimization
- ③ Balancing Fidelity and Diversity
- ④ Scalability to Large-scale and Diverse Modalities
- ⑤ Ethical and Privacy Considerations